



```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from scipy.stats import pointbiserialr, spearmanr
from sklearn.feature_selection import SelectKBest
from sklearn.model_selection import cross_val_score
from sklearn.tree import DecisionTreeClassifier
```

```
df=pd.read_csv("adult.csv")
df
```

	age	workclass	fnlwgt	education	education.num	marital.status	occupatio
0	90	?	77053	HS-grad	9	Widowed	
1	82	Private	132870	HS-grad	9	Widowed	Exec manageri
2	66	?	186061	Some- college	10	Widowed	
3	54	Private	140359	7th-8th	4	Divorced	Machine op-inspe
4	41	Private	264663	Some- college	10	Separated	Pro specialt
...	...	...	...	...	...	...	.
32556	22	Private	310152	Some- college	10	Never-married	Protective ser
32557	27	Private	257302	Assoc- acdm	12	Married-civ- spouse	Tech suppo
32558	40	Private	154374	HS-grad	9	Married-civ- spouse	Machine op-inspe
32559	58	Private	151910	HS-grad	9	Widowed	Adm clerica
32560	22	Private	201490	HS-grad	9	Never-married	Adm clerica

32561 rows × 15 columns



Data Analysis

```
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   32561 non-null  int64
1   workclass             32561 non-null  object
2   fnlwgt                32561 non-null  int64
3   education             32561 non-null  object
4   education.num         32561 non-null  int64
5   marital.status        32561 non-null  object
6   occupation            32561 non-null  object
7   relationship          32561 non-null  object
8   race                  32561 non-null  object
9   sex                   32561 non-null  object
10  capital.gain           32561 non-null  int64
11  capital.loss           32561 non-null  int64
12  hours.per.week        32561 non-null  int64
13  native.country        32561 non-null  object
14  income                 32561 non-null  object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB
```

```
df.describe()
```

	age	fnlwgt	education.num	capital.gain	capital.loss	hours
count	32561.000000	3.256100e+04	32561.000000	32561.000000	32561.000000	32561.000000
mean	38.581647	1.897784e+05	10.080679	1077.648844	87.303830	40.921004
std	13.640433	1.055500e+05	2.572720	7385.292085	402.960219	13.959118
min	17.000000	1.228500e+04	1.000000	0.000000	0.000000	0.000000
25%	28.000000	1.178270e+05	9.000000	0.000000	0.000000	0.000000
50%	37.000000	1.783560e+05	10.000000	0.000000	0.000000	0.000000
75%	48.000000	2.370510e+05	12.000000	0.000000	0.000000	0.000000
max	90.000000	1.484705e+06	16.000000	99999.000000	4356.000000	99.000000

df.head()

	age	workclass	fnlwgt	education	education.num	marital.status	occupation	relationship
0	90	?	77053	HS-grad	9	Widowed	?	Widowed
1	82	Private	132870	HS-grad	9	Widowed	Exec-managerial	Widowed
2	66	?	186061	Some-college	10	Widowed	?	Widowed
3	54	Private	140359	7th-8th	4	Divorced	Machine-op-inspct	Divorced
4	41	Private	264663	Some-college	10	Separated	Prof-specialty	Separated

#checking total missing values in the whole dataset of attributes

```
df_missing = (df=='?').sum()
```

```
df_missing
```

```
age          0
workclass    1836
fnlwgt       0
education    0
education.num 0
marital.status 0
occupation   1843
relationship  0
race         0
sex          0
capital.gain  0
capital.loss  0
hours.per.week 0
native.country 583
income       0
dtype: int64
```

```
percent_missing = (df=='?').sum() * 100/len(df)
```

```
percent_missing
```


```
age          0.000000
workclass    5.638647
fnlwgt       0.000000
education    0.000000
education.num 0.000000
marital.status 0.000000
occupation   5.660146
relationship  0.000000
race         0.000000
sex          0.000000
capital.gain  0.000000
capital.loss  0.000000
hours.per.week 0.000000
native.country 1.790486
income       0.000000
dtype: float64
```

```
df.apply(lambda x: x != '?',axis=1).sum() #rows not containing "?"
```

```
age          32561
workclass    30725
fnlwgt       32561
education    32561
education.num 32561
marital.status 32561
occupation   30718
relationship 32561
race         32561
sex          32561
capital.gain 32561
capital.loss 32561
hours.per.week 32561
native.country 31978
income       32561
dtype: int64
```

```
df = df[df['workclass'] != '?'] #dropping missing value rows
df.head()
```

	age	workclass	fnlwgt	education	education.num	marital.status	occupation	native.country
1	82	Private	132870	HS-grad	9	Widowed	Exec-managerial	
3	54	Private	140359	7th-8th	4	Divorced	Machine-op-inspct	
4	41	Private	264663	Some-college	10	Separated	Prof-specialty	
5	34	Private	216864	HS-grad	9	Divorced	Other-service	
6	38	Private	150601	10th	6	Separated	Adm-clerical	



```
df_categorical = df.select_dtypes(include=['object'])
df_categorical.apply(lambda x: x == '?',axis=1).sum()
```

```
workclass      0
education      0
marital.status  0
occupation      7
relationship    0
race           0
sex            0
native.country 556
income         0
dtype: int64
```

```
df = df[df['occupation'] != '?']
df = df[df['native.country'] != '?']
#final check for null values
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 30162 entries, 1 to 32560
Data columns (total 15 columns):
#   Column          Non-Null Count  Dtype
---  -
0   age             30162 non-null  int64
1   workclass       30162 non-null  object
2   fnlwgt          30162 non-null  int64
3   education       30162 non-null  object
4   education.num   30162 non-null  int64
5   marital.status  30162 non-null  object
6   occupation      30162 non-null  object
7   relationship     30162 non-null  object
8   race            30162 non-null  object
9   sex             30162 non-null  object
10  capital.gain     30162 non-null  int64
11  capital.loss     30162 non-null  int64
12  hours.per.week   30162 non-null  int64
13  native.country   30162 non-null  object
```

```
14 income          30162 non-null object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB
```

Data preprocessing

```
from sklearn import preprocessing
```

```
# encoding categorical variables using label Encoder
df_categorical = df.select_dtypes(include=['object'])
df_categorical.head()
```

	workclass	education	marital.status	occupation	relationship	race	sex	r
1	Private	HS-grad	Widowed	Exec-manual	Not-in-family	White	Female	
3	Private	7th-8th	Divorced	Machine-op-inspct	Unmarried	White	Female	
4	Private	Some-college	Separated	Prof-specialty	Own-child	White	Female	

```
le = preprocessing.LabelEncoder()
df_categorical = df_categorical.apply(le.fit_transform)
df_categorical.head()
```

	workclass	education	marital.status	occupation	relationship	race	sex	nati
1	2	11	6	3	1	4	0	
3	2	5	0	6	4	4	0	
4	2	15	5	9	3	4	0	
5	2	11	0	7	4	4	0	
6	2	0	5	0	4	4	1	

```
#dropping duplicate columns which had categorical values
df = df.drop(df_categorical.columns,axis=1)
df = pd.concat([df,df_categorical],axis=1)
df.head()
```

	age	fnlwgt	education.num	capital.gain	capital.loss	hours.per.week	workcla
1	82	132870	9	0	4356	18	
3	54	140359	4	0	3900	40	
4	41	264663	10	0	3900	40	
5	34	216864	9	0	3770	45	
6	38	150601	6	0	3770	40	



```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 30162 entries, 1 to 32560
Data columns (total 15 columns):
#   Column              Non-Null Count  Dtype
---  -
0   age                  30162 non-null  int64
1   fnlwgt               30162 non-null  int64
2   education.num        30162 non-null  int64
3   capital.gain         30162 non-null  int64
4   capital.loss         30162 non-null  int64
5   hours.per.week       30162 non-null  int64
6   workclass            30162 non-null  int64
7   education            30162 non-null  int64
8   marital.status       30162 non-null  int64
9   occupation           30162 non-null  int64
10  relationship         30162 non-null  int64
```

```

11 race                30162 non-null int64
12 sex                 30162 non-null int64
13 native.country     30162 non-null int64
14 income              30162 non-null int64
dtypes: int64(15)
memory usage: 3.7 MB

```

```

# converting target variable income to categorical
df['income'] = df['income'].astype('category')
df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 30725 entries, 1 to 32560
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   30725 non-null  int64
1   workclass              30725 non-null  object
2   fnlwgt                 30725 non-null  int64
3   education              30725 non-null  object
4   education.num          30725 non-null  int64
5   marital.status         30725 non-null  object
6   occupation             30725 non-null  object
7   relationship           30725 non-null  object
8   race                   30725 non-null  object
9   sex                   30725 non-null  object
10  capital.gain           30725 non-null  int64
11  capital.loss           30725 non-null  int64
12  hours.per.week         30725 non-null  int64
13  native.country         30725 non-null  object
14  income                 30725 non-null  category
dtypes: category(1), int64(6), object(8)
memory usage: 3.5+ MB
<ipython-input-56-f55635cbcb0>:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df['income'] = df['income'].astype('category')
```

```

from sklearn.model_selection import train_test_split
# Putting independent variables/features to X
X = df.drop('income',axis=1)

# Putting response/dependent variable/feature to y
y = df['income']
X.head(3)

```

	age	workclass	fnlwgt	education	education.num	marital.status	occupation	income
1	82	Private	132870	HS-grad	9	Widowed	Exec-managerial	<=50K
3	54	Private	140359	7th-8th	4	Divorced	Machine-op-inspct	<=50K
4	41	Private	264663	Some-college	10	Separated	Prof-specialty	<=50K



```
y.head(3)
```

```

1    <=50K
3    <=50K
4    <=50K
Name: income, dtype: category
Categories (2, object): ['<=50K', '>50K']

```

splitting the values in test & train datasets

```
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.30,random_state=99)
```

```
X_train.head()
```

	age	workclass	fnlwgt	education	education.num	marital.status	occupatio
32493	51	Private	177669	Some-college	10	Married-civ-spouse	Sale
4725	20	Private	27337	HS-grad	9	Never-married	Handlers/cleaner
6380	38	Private	258761	HS-grad	9	Divorced	Executive/managerial
10782	43	Private	120277	HS-grad	9	Married-civ-spouse	Craft-repair
29936	21	Private	199419	Some-college	10	Never-married	Administrative/clerical



```
from sklearn.tree import DecisionTreeClassifier
```

```
# max_depth which is 5 so that we can plot and read the tree.
dt_default = DecisionTreeClassifier(max_depth=5)
dt_default.fit(X_train,y_train)
```

```
-----
ValueError                                Traceback (most recent call last)
<ipython-input-60-6cfc0f7d38c9> in <cell line: 5>()
      3 # max_depth which is 5 so that we can plot and read the tree.
      4 dt_default = DecisionTreeClassifier(max_depth=5)
----> 5 dt_default.fit(X_train,y_train)

----- 5 frames -----
/usr/local/lib/python3.10/dist-packages/pandas/core/generic.py in __array__(self, dtype)
    2068
    2069     def __array__(self, dtype: npt.DTypeLike | None = None) -> np.ndarray:
-> 2070         return np.asarray(self._values, dtype=dtype)
    2071
    2072     def __array_wrap__(

ValueError: could not convert string to float: 'Private'
```

SEARCH STACK OVERFLOW

```
from sklearn.metrics import classification_report,confusion_matrix,accuracy_score
```

```
y_pred_default = dt_default.predict(X_test)
print(classification_report(y_test,y_pred_default)) #classifier
```

```
-----
ValueError                                Traceback (most recent call last)
<ipython-input-61-7c9dee1e4662> in <cell line: 3>()
      1 from sklearn.metrics import
      2 classification_report,confusion_matrix,accuracy_score
----> 3 y_pred_default = dt_default.predict(X_test)
      4 print(classification_report(y_test,y_pred_default)) #classifier

----- 5 frames -----
/usr/local/lib/python3.10/dist-packages/pandas/core/generic.py in __array__(self, dtype)
    2068
    2069     def __array__(self, dtype: npt.DTypeLike | None = None) -> np.ndarray:
-> 2070         return np.asarray(self._values, dtype=dtype)
    2071
    2072     def __array_wrap__(

ValueError: could not convert string to float: 'Private'
```

SEARCH STACK OVERFLOW

```
print(confusion_matrix(y_test,y_pred_default))
print(accuracy_score(y_test,y_pred_default))
```

```
[[6553  314]
 [1039 1143]]
0.8504807161012267
```

```
# Putting features
features = list(df.columns[1:])
features
```

```
dot_data = StringIO()
export_graphviz(dt_default, out_file=dot_data, feature_names=features, filled=True, rounded=True)

graph = pydotplus.graph_from_dot_data(dot_data.getvalue())
Image(graph.create_png())
```

